

# Bike Sharing - Hypothesis Testing

May 21, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import ttest_ind, chi2_contingency, f_oneway, levene
from statsmodels.graphics.gofplots import qqplot
```

---

## 1 Objective :

A bike sharing company has recently faced a significant dip in its revenue and we would like to understand why. Especially as data scientists, determine the factors which are good predictors of demand with a certain level of confidence. Further on the basis of these strong predictors of demand we also need to give recommendations to the company.

## 2 Insights :

- Clear, working winter day is the most common day reported in the data
- The count of casual, registered and total rental bikes per hour shows a somewhat poisson distribution
- Temp, feeling temp, humidity and windspeed are also somewhat normally distributed centered around 20, 25, 60 and 13 respectively
- Insights from visual EDA :
  - Most bikes seems to be rented in fall and the least in spring
  - There seems to be slightly more number of bikes rented on a holiday as compared to no holiday (very small difference)
  - There seems to be slightly lesser number of bikes rented on a working day (very small difference)
  - Most bikes are rented on clear days and the least on rainy days
- We cannot infer that working day has an effect on number of bikes rented (t-test pvalue is 22.6%)

- We can infer that weather has a significant effect on number of bikes rented (f-test pvalue is ~0%)
- We can infer that clear weather is linked to increase in number of bikes rented (Welch t-test pvalue ~0%)
- We can infer that season has a significant effect on number of bikes rented (f-test pvalue is ~0%)
- We can infer that summer and fall season is linked to increase in number of bikes rented (Welch t-test pvalue ~0%)
- We can infer that season has a significant effect on the weather (chisquare-test pvalue is ~0%)

### 3 Recommendations :

- The bike sharing company should start extensive marketing campaigns at the start of summer season and run them all the way through fall season.
- During these 2 seasons they should increase their supply of cycles, and install more pick up/drop locations to make sure that the demand is fully met.
- They should also introduce Miles Points (like flights) during these seasons to build customer loyalty and retain them for the next year as well.

## 4 EDA

### 4.1 Non Visual

```
[2]: df = pd.read_csv('bike_sharing.csv')
```

```
[3]: df.shape
```

```
[3]: (10886, 12)
```

```
[4]: df.head()
```

```
[4]:
```

	datetime	season	holiday	workingday	weather	temp	atemp	\
0	2011-01-01 00:00:00	1	0	0	1	9.84	14.395	
1	2011-01-01 01:00:00	1	0	0	1	9.02	13.635	
2	2011-01-01 02:00:00	1	0	0	1	9.02	13.635	
3	2011-01-01 03:00:00	1	0	0	1	9.84	14.395	
4	2011-01-01 04:00:00	1	0	0	1	9.84	14.395	

	humidity	windspeed	casual	registered	count
0	81	0.0	3	13	16
1	80	0.0	8	32	40
2	80	0.0	5	27	32

3	75	0.0	3	10	13
4	75	0.0	0	1	1

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10886 entries, 0 to 10885
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   datetime         10886 non-null  object
1   season           10886 non-null  int64
2   holiday          10886 non-null  int64
3   workingday       10886 non-null  int64
4   weather          10886 non-null  int64
5   temp             10886 non-null  float64
6   atemp            10886 non-null  float64
7   humidity         10886 non-null  int64
8   windspeed        10886 non-null  float64
9   casual           10886 non-null  int64
10  registered       10886 non-null  int64
11  count            10886 non-null  int64
dtypes: float64(3), int64(8), object(1)
memory usage: 1020.7+ KB
```

- No nulls found in data

```
[6]: #Changing dtype:
df['datetime'] = pd.to_datetime(df['datetime'], format='%Y-%m-%d %H:%M:%S')

#Coverting to categorical values for easier interpretation of data
df['season'] = df['season'].replace({1:'spring',2:'summer',3:'fall',4:'winter'})
df['holiday'] = df['holiday'].replace({1:'holiday',0:'no holiday'})
df['workingday'] = df['workingday'].replace({1:'workingday',0:'no workingday'})
df['weather'] = df['weather'].replace({1:'clear',2:'mist',3:'light rain/snow',4:
    ↳'heavy rain/snow'})
```

```
[7]: df.describe(datetime_is_numeric=True)
```

```
[7]:
```

	datetime	temp	atemp	humidity \
count	10886	10886.00000	10886.000000	10886.000000
mean	2011-12-27 05:56:22.399411968	20.23086	23.655084	61.886460
min	2011-01-01 00:00:00	0.82000	0.760000	0.000000
25%	2011-07-02 07:15:00	13.94000	16.665000	47.000000
50%	2012-01-01 20:30:00	20.50000	24.240000	62.000000
75%	2012-07-01 12:45:00	26.24000	31.060000	77.000000
max	2012-12-19 23:00:00	41.00000	45.455000	100.000000
std	NaN	7.79159	8.474601	19.245033

	windspeed	casual	registered	count
count	10886.000000	10886.000000	10886.000000	10886.000000
mean	12.799395	36.021955	155.552177	191.574132
min	0.000000	0.000000	0.000000	1.000000
25%	7.001500	4.000000	36.000000	42.000000
50%	12.998000	17.000000	118.000000	145.000000
75%	16.997900	49.000000	222.000000	284.000000
max	56.996900	367.000000	886.000000	977.000000
std	8.164537	49.960477	151.039033	181.144454

```
[8]: df.describe(include='O')
```

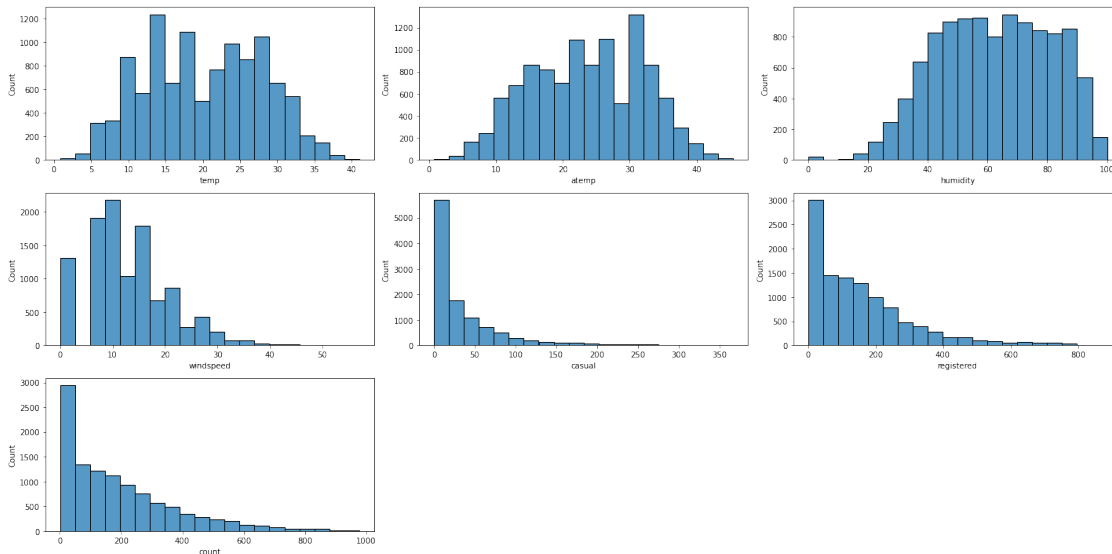
```
[8]:
      season  holiday  workingday  weather
count    10886     10886      10886    10886
unique         4         2         2         4
top    winter  no holiday  workingday  clear
freq     2734     10575       7412     7192
```

- Clear, working winter day is the most common day reported in the data

## 4.2 Visual EDA

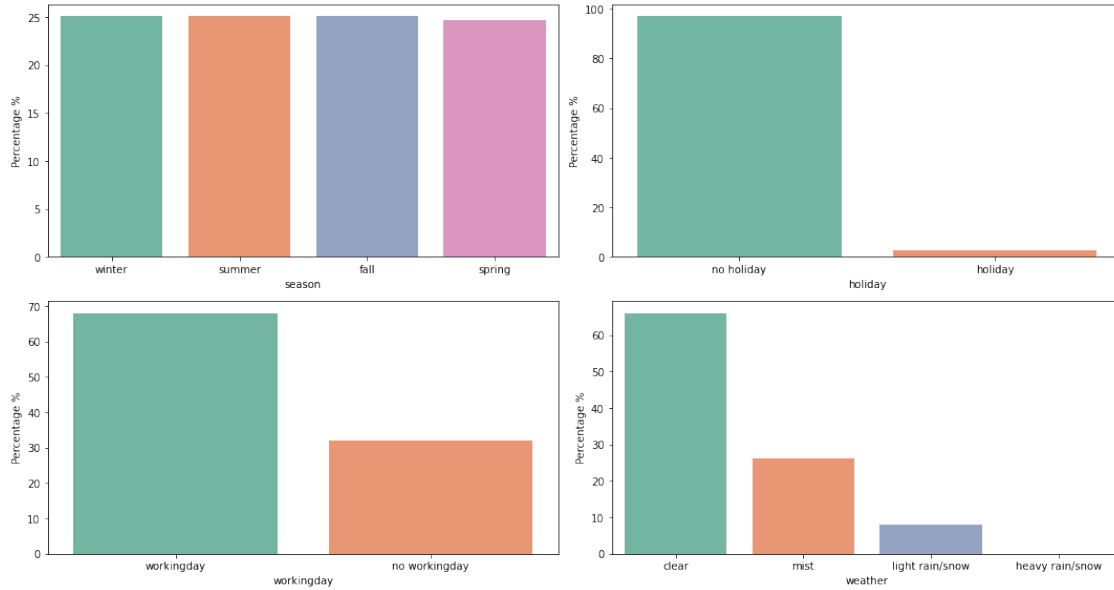
### 4.2.1 Univariate

```
[10]: plt.rcParams['figure.figsize'] = (20,10)
plt.rcParams['figure.autolayout'] = True
continous_variables = ['temp', 'atemp', 'humidity', 'windspeed', 'casual', 'registered', 'count']
for e, c in enumerate(continous_variables):
    ax = plt.subplot(3,3,e+1)
    sns.histplot(df[c],ax=ax, bins=20)
```



- The count of casual, registered and total rental bikes per hour shows a somewhat poisson distribution
- Temp, feeling temp, humidity and windspeed are also somewhat normally distributed centered around 20, 25, 60 and 13 respectively

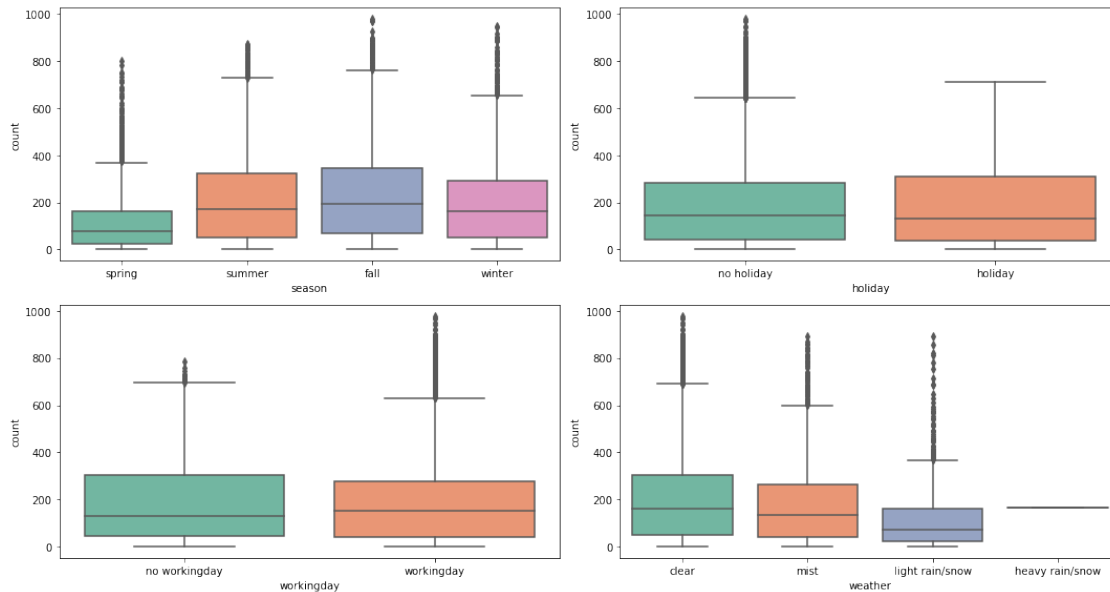
```
[11]: plt.rcParams['figure.figsize'] = (15,8)
categorical_variables = ['season', 'holiday', 'workingday', 'weather']
for e,c in enumerate(categorical_variables):
    ax = plt.subplot(2,2,e+1)
    data = df[c].value_counts(1) * 100
    sns.barplot(x=data.index, y=data.values, ax=ax, palette='Set2')
    ax.set_ylabel('Percentage %')
    ax.set_xlabel(c)
```



- We have equal data for all the 4 seasons
- We have less than 10% data from holidays (Because there are very few holidays in a year)
- Around 65% of the days are working in a year (Based on holidays + weekends)
- Around 65% of the days had clear weather, and almost 0% days had heavy rain

## 4.2.2 Bivariate

```
[12]: fig, ax = plt.subplots(2,2,figsize=(15,8))
sns.boxplot(x='season', y='count', data=df, ax=ax[0][0], palette='Set2')
sns.boxplot(x='holiday', y='count', data=df, ax=ax[0][1], palette='Set2')
sns.boxplot(x='workingday', y='count', data=df, ax=ax[1][0], palette='Set2')
sns.boxplot(x='weather', y='count', data=df, ax=ax[1][1], palette='Set2')
print()
```



- From the above boxplots we can say :
  - Most bikes seems to be rented in fall and the least in spring
  - There seems to be slightly more number of bikes rented on a holiday as compared to no holiday (very small difference)
  - There seems to be slightly lesser number of bikes rented on a working day (very small difference)
  - Most bikes are rented on clear days and the least on rainy days

## 5 Hypothesis Testing

```
[13]: df_copy = pd.read_csv('bike_sharing.csv')
```

### 5.1 Does Working Day effects bike rentals?

- $H_0$  : Working day has no effect on number of bikes rented
- $H_a$  : Working day has an effect on number of bikes rented
- Significance level : 5%
- Test Statistic : t-test score (2 Sample t-test)

```
[14]: bikes_on_working_day = df_copy.loc[df_copy['workingday'] == 1, 'count'].values
      bikes_on_non_working_day = df_copy.loc[df_copy['workingday'] == 0, 'count'].
      ↪ values
```

### 5.1.1 Checking ttest assumptions

```
[15]: # Equal Variance : Both samples need to have equal variance otherwise we will
      ↪ do a Welch t-test
      # Null hypothesis : There is no difference of standard deviation between the
      ↪ samples
      # Alternate hypothesis : There is a difference of standard deviation between
      ↪ the samples
      # We can do a levene test to check this at a significance level of 5%

      levene(bikes_on_working_day, bikes_on_non_working_day)
```

```
[15]: LeveneResult(statistic=0.004972848886504472, pvalue=0.9437823280916695)
```

- The levene test gives a pvalue of 94% which is much higher than our significance level
- Hence we cannot reject the null hypothesis
- Hence we cannot say that these distributions have different standard deviations

```
[16]: # We will do a 2 sample t-test to check if these two are independent or not
      ttest_ind(bikes_on_working_day, bikes_on_non_working_day)
```

```
[16]: Ttest_indResult(statistic=1.2096277376026694, pvalue=0.22644804226361348)
```

- We get a t-test score of 1.2
- Our pvalue is 22.6% which is much higher than our significance level of 5%
- We cannot reject the H0
- We cannot infer that working day has an effect on number of bikes rented

---

## 5.2 Does Weather effects bike rentals?

```
[17]: df_copy['weather'].value_counts()
```

```
[17]: 1    7192
      2    2834
      3     859
      4      1
      Name: weather, dtype: int64
```

- There is only 1 data point for weather type 4 ie (Heavy Rain + Ice Pallets + Thunderstorm + Mist, Snow + Fog)
- We will do ANOVA on the rest 3 weather types only
- H0 : Weather has no effect on number of bikes rented
- Ha : Weather has an effect on number of bikes rented
- Significance level : 5%
- Test Statistic : f-statistic (ANOVA)

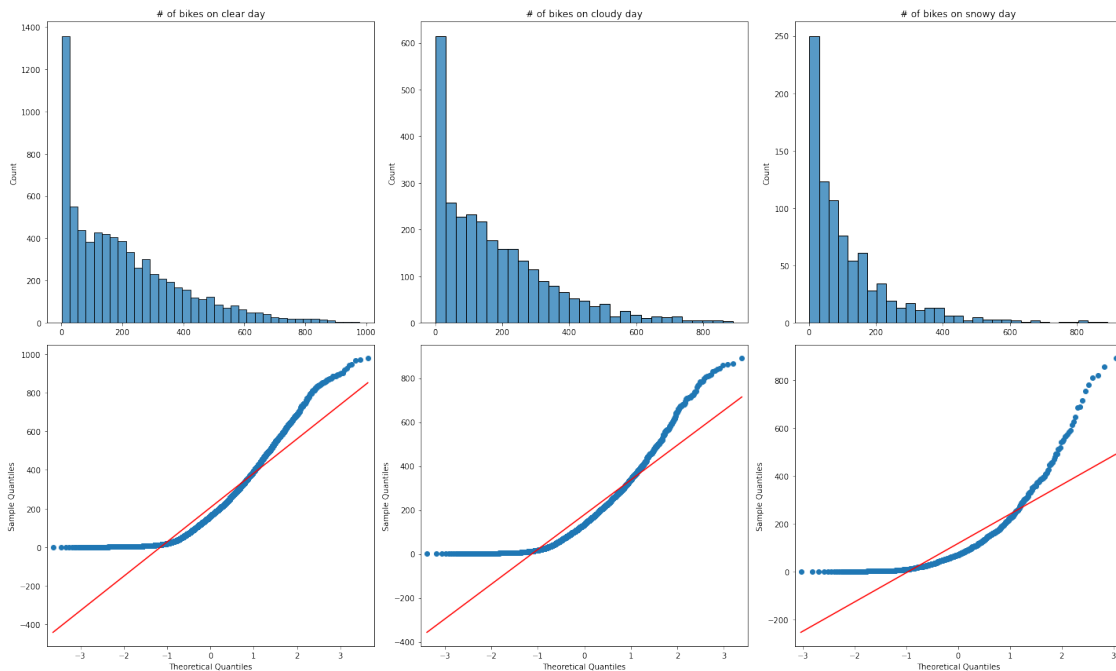


```
[18]: bikes_on_weather_type_1 = df_copy.loc[df_copy['weather'] == 1, 'count'].values
      bikes_on_weather_type_2 = df_copy.loc[df_copy['weather'] == 2, 'count'].values
      bikes_on_weather_type_3 = df_copy.loc[df_copy['weather'] == 3, 'count'].values
```

### 5.2.1 Checking ANOVA assumptions

```
[19]: # Normal Assumption : The sample distributions should be normally distributed
fig, ax = plt.subplots(2, 3, figsize=(20, 12))
sns.histplot(bikes_on_weather_type_1, ax=ax[0][0])
sns.histplot(bikes_on_weather_type_2, ax=ax[0][1])
sns.histplot(bikes_on_weather_type_3, ax=ax[0][2])
ax[0][0].set_title('# of bikes on clear day')
ax[0][1].set_title('# of bikes on cloudy day')
ax[0][2].set_title('# of bikes on snowy day')

qqplot(bikes_on_weather_type_1, line='r', ax=ax[1][0])
qqplot(bikes_on_weather_type_2, line='r', ax=ax[1][1])
qqplot(bikes_on_weather_type_3, line='r', ax=ax[1][2])
plt.show()
```



- We can see from above graphs that none of the distributions are normal

```
[20]: # Equal Variance Assumption : The sample distribution should have equal variance
```

```
# Null hypothesis : There is no difference of standard deviation between the
↳ samples
# Alternate hypothesis : There is a difference of standard deviation between
↳ the samples
# We can do a levene test to check this at a significance level of 5%

levene(bikes_on_weather_type_1, bikes_on_weather_type_2,
↳ bikes_on_weather_type_3)
```

[20]: LeveneResult(statistic=81.67574924435011, pvalue=6.198278710731511e-36)

- The levene test gives a pvalue of ~0% which is much less than our significance level
- Hence we can reject the null hypothesis
- Hence we can say that these distributions have different standard deviations

```
[21]: # We now do a one way ANOVA to check the difference of means between the samples
f_oneway(bikes_on_weather_type_1, bikes_on_weather_type_2,
↳ bikes_on_weather_type_3)
```

[21]: F\_onewayResult(statistic=98.28356881946706, pvalue=4.976448509904196e-43)

- We get a fstat of 98.3
- Our pvalue is ~0% which is much lower than our significance level of 5%
- We reject the H0
- We can infer that weather has a significant effect on number of bikes rented

### 5.3 Are there more bikes rented on clear days?

- H0 : Clear weather has no effect on number of bikes rented
- Ha : Clear weather increases the number of bikes rented
- Significance level : 5%
- Test Statistic : t-test score (2 sample t-test)

```
[22]: bikes_on_weather_type_2_and_3 = np.concatenate((bikes_on_weather_type_2,
↳ bikes_on_weather_type_3))
```

```
[23]: # Checking ttest assumptions
# Equal Variance Assumption : The sample distribution should have equal variance
# Null hypothesis : There is no difference of standard deviation between the
↳ samples
# Alternate hypothesis : There is a difference of standard deviation between
↳ the samples
# We can do a levene test to check this at a significance level of 5%

levene(bikes_on_weather_type_1, bikes_on_weather_type_2_and_3)
```

[23]: LeveneResult(statistic=87.13479803756756, pvalue=1.210545470625517e-20)

- The levene test gives a pvalue of ~0% which is much less than our significance level
- Hence we can reject the null hypothesis
- Hence we can say that these distributions have different standard deviations
- Hence we will do a Welch t-test instead here

```
[24]: # We will do a 2 sample Welch t-test to check if these two are independent or
      ↪not
      ttest_ind(bikes_on_weather_type_1, bikes_on_weather_type_2_and_3,
      ↪alternative='greater', equal_var=False)
```

```
[24]: Ttest_indResult(statistic=11.5340767638041, pvalue=7.540763883724984e-31)
```

- We get a t-test score of 11
- Our pvalue is ~0% which is much lower than our significance level of 5%
- We can reject the H0
- We can infer that clear weather is linked to increase in number of bikes rented

## 5.4 Does season effects bike rentals?

```
[25]: df_copy['season'].value_counts()
```

```
[25]: 4    2734
      2    2733
      3    2733
      1    2686
      Name: season, dtype: int64
```

- H0 : Season has no effect on number of bikes rented
- Ha : Season has an effect on number of bikes rented
- Significance level : 5%
- Test Statistic : f-statistic (ANOVA)

```
[26]: bikes_on_season_type_1 = df_copy.loc[df_copy['season'] == 1, 'count'].values
      bikes_on_season_type_2 = df_copy.loc[df_copy['season'] == 2, 'count'].values
      bikes_on_season_type_3 = df_copy.loc[df_copy['season'] == 3, 'count'].values
      bikes_on_season_type_4 = df_copy.loc[df_copy['season'] == 4, 'count'].values
```

### 5.4.1 Checking ANOVA assumptions

```
[27]: # Normal Assumption : The sample distributions should be normally distributed
      fig, ax = plt.subplots(2, 4, figsize=(25, 12))
      sns.histplot(bikes_on_season_type_1, ax=ax[0][0])
      sns.histplot(bikes_on_season_type_2, ax=ax[0][1])
      sns.histplot(bikes_on_season_type_3, ax=ax[0][2])
```

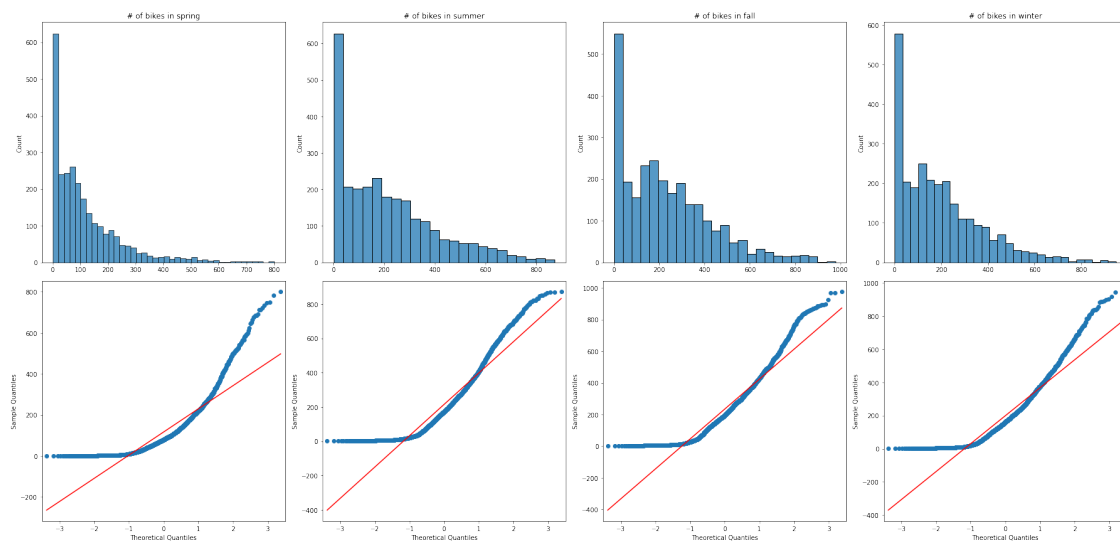
```

sns.histplot(bikes_on_season_type_4,ax=ax[0][3])
ax[0][0].set_title('# of bikes in spring')
ax[0][1].set_title('# of bikes in summer')
ax[0][2].set_title('# of bikes in fall')
ax[0][3].set_title('# of bikes in winter')

qqplot(bikes_on_season_type_1,line='r',ax=ax[1][0])
qqplot(bikes_on_season_type_2,line='r',ax=ax[1][1])
qqplot(bikes_on_season_type_3,line='r',ax=ax[1][2])
qqplot(bikes_on_season_type_4,line='r',ax=ax[1][3])

plt.show()

```



- We can see from above graphs that none of the distributions are normal

```

[28]: # Equal Variance Assumption : The sample distribution should have equal variance
# Null hypothesis : There is no difference of standard deviation between the
      ↪ samples
# Alternate hypothesis : There is a difference of standard deviation between
      ↪ the samples
# We can do a levene test to check this at a significance level of 5%

levene(bikes_on_season_type_1, bikes_on_season_type_2, bikes_on_season_type_3,
      ↪ bikes_on_season_type_4)

```

```

[28]: LeveneResult(statistic=187.7706624026276, pvalue=1.0147116860043298e-118)

```

- The levene test gives a pvalue of ~0% which is much less than our significance level
- Hence we can reject the null hypothesis
- Hence we can say that these distributions have different standard deviations

```
[29]: # We will now do one way ANOVA to check difference between means
f_oneway(bikes_on_season_type_1, bikes_on_season_type_2,
↳bikes_on_season_type_3, bikes_on_season_type_4)
```

```
[29]: F_onewayResult(statistic=236.94671081032106, pvalue=6.164843386499654e-149)
```

- We get a fstat of 236.9
- Our pvalue is ~0% which is much lower than our significance level of 5%
- We reject the H0
- We can infer that season has a significant effect on number of bikes rented

## 5.5 Are there more bikes rented in summer and fall?

- H0 : Summer and Fall season has no effect on number of bikes rented
- Ha : Summer and Fall season increases the number of bikes rented
- Significance level : 5%
- Test Statistic : t-test score (2 sample t-test)

```
[30]: bikes_on_season_type_2_and_3 = np.concatenate((bikes_on_season_type_2,
↳bikes_on_season_type_3))
bikes_on_season_type_1_and_4 = np.concatenate((bikes_on_season_type_1,
↳bikes_on_season_type_4))
```

```
[31]: # Checking ttest assumptions
# Equal Variance Assumption : The sample distribution should have equal variance
# Null hypothesis : There is no difference of standard deviation between the
↳samples
# Alternate hypothesis : There is a difference of standard deviation between
↳the samples
# We can do a levene test to check this at a significance level of 5%

levene(bikes_on_season_type_2_and_3, bikes_on_season_type_1_and_4)
```

```
[31]: LeveneResult(statistic=237.21717358575066, pvalue=5.74764555857786e-53)
```

- The levene test gives a pvalue of ~0% which is much less than our significance level
- Hence we can reject the null hypothesis
- Hence we can say that these distributions have different standard deviations
- Hence we will do a Welch t-test instead here

```
[32]: # We will do a 2 sample Welch t-test to check if these two are independent or
↳not
ttest_ind(bikes_on_season_type_2_and_3, bikes_on_season_type_1_and_4,
↳alternative='greater', equal_var=False)
```

```
[32]: Ttest_indResult(statistic=19.589523172181725, pvalue=2.9605129625749147e-84)
```

- We get a t-test score of 19.5
  - Our pvalue is ~0% which is much lower than our significance level of 5%
  - We can reject the H0
  - We can infer that summer and fall season is linked to increase in number of bikes rented
- 

## 5.6 Is weather dependent on season?

```
[33]: df_copy['weather'].value_counts()
```

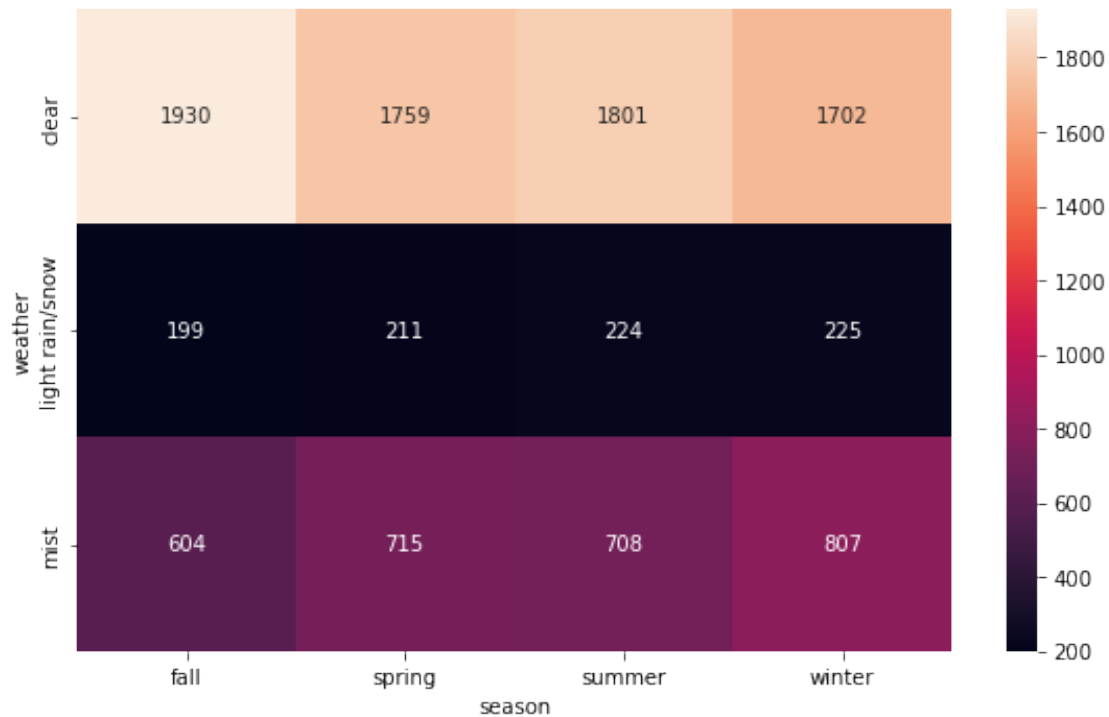
```
[33]: 1    7192
      2    2834
      3     859
      4      1
      Name: weather, dtype: int64
```

- Again dropping the weather type 4 before doing the test
- In chi2 test we also require atleast 5 samples in each category
- H0 : Season has no effect on the weather
- Ha : Season has an effect on the weather
- Significance level : 5%
- Test Statistic : chi-statistic

```
[34]: df_copy = df_copy[df_copy['weather'] != 4]
      df = df[df['weather'] != 'heavy rain/snow']
```

```
[35]: #Lets look at a crosstab between weather and season
      plt.rcParams['figure.figsize'] = (8,5)
      sns.heatmap(pd.crosstab(df['weather'], df['season']),annot=True,fmt='d')
```

```
[35]: <AxesSubplot:xlabel='season', ylabel='weather'>
```



```
[36]: # We will do a chi square test to see if these two features are dependent or not
chi2_contingency(pd.crosstab(df_copy['weather'], df_copy['season']))
```

```
[36]: (46.101457310732485,
2.8260014509929403e-08,
6,
array([[1774.04869086, 1805.76352779, 1805.76352779, 1806.42425356],
      [ 699.06201194,  711.55920992,  711.55920992,  711.81956821],
      [ 211.8892972 ,  215.67726229,  215.67726229,  215.75617823]]))
```

- We get a chisquare-stat of 46.1
- Our pvalue is ~0% which is much lower than our significance level of 5%
- We reject the H0
- We can infer that season has a significant effect on the weather